

**A Deep Learning-Based Perception-Driven Vectorization
Approach For Semi-Structured Imagery**

A Special Problem by

Sean Francis N. Ballais

2015-04562

BS CS

**Presented to the Faculty of the
Division of Natural Sciences and Mathematics**

**In Partial Fulfillment of the Requirements
For the Degree of
Bachelor of Science in Computer Science**

**University of the Philippines Visayas
TACLOBAN COLLEGE
Tacloban City**

Month Year

This special problem, entitled “**A DEEP LEARNING-BASED PERCEPTION-DRIVEN VECTORIZATION APPROACH FOR SEMI-STRUCTURED IMAGERY**”, prepared and submitted by **SEAN FRANCIS N. BALLAIS**, in partial fulfillment of the requirements for the degree of **BACHELOR OF SCIENCE IN COMPUTER SCIENCE** is hereby accepted.

PROF. VICTOR M. ROMERO II
Special Problem Adviser

Accepted as partial fulfillment of the requirements for the degree of **BACHELOR OF SCIENCE IN COMPUTER SCIENCE**.

DR. EULITO V. CASAS JR.
Chair, DNSM

Acknowledgements

First of all I would like to thank the Lord for his guidance during the course of my research and for making this thesis possible. I would like to thank my family who served as my inspiration, and never failed to support me all throughout my studies. Thanks to ...

Abstract

Although the abstract is the first thing that appears in the thesis, it is best written last after you have written your conclusion. It should contain spell out your thesis problem and describe your solution clearly.

Make sure your abstract fits in one page !

Table of Contents

Acknowledgements	iii
Abstract	iv
Table of Contents	vi
List of Tables	vii
List of Figures	viii
1 Introduction	1
1.1 The University Course Timetabling Problem	2
1.2 Algorithms Used In Approaches For Solving the University Course Timetabling Problem	3
Great Deluge	3
Genetic Algorithms	4
2 Review of Related Literature	7
2.1 Great Deluge-Based Works	7
2.2 Machine Learning-Based Works	11
2.3 Genetic Algorithm-Based Works	14
3 Statement of the Problem	21
4 Objectives	22
5 Proposed Methodology	23
5.1 Image Preprocessing	23
Corner Detection	23
Region Segmentation	24

	vi
5.2 Curve Approximation	25
Point Sequence Generation	25
Point Sequence Curve Approximation	26
5.3 Region Colouring and Merging	33
6 Describing How You Validated Your Approach.	35
7 Stating Your Results and Drawing Insights From Them.	36
8 Summarizing Your Thesis and Drawing Your Conclusions.	37
A What should be in the Appendix	38
Bibliography	39

List of Tables

List of Figures

5.1	The network architecture of the Point Parametrization Network (PPN)	28
5.2	The network architecture of the Knot Selection Network (PPN) . . .	31

Chapter 1

Introduction

A common activity in universities and all academic institutions is course timetabling. This activity is a necessity in academic institutions as this is where classes are scheduled for the incoming semester. In some instances, even entire courses are scheduled for at least four years [3]. Performing timetabling manually is tedious and schedules conflicts may not immediately be determined [17]. As such, automating this procedure will be beneficial to academic institutions. This problem of automating course scheduling is actually formally known as the University Course Timetabling Problem (UTP) [21][20]. Many works have already proposed different approaches to automate the timetabling process. Some of the algorithms that were utilized by proposed approaches include great deluge [?][16], machine learning [17], and genetic algorithms [4][19][15]. Their results produce optimal timetables, or near to the optimal, at least for the case they were originally intended for. Applying their approaches to other institutions *may* require some tweaking to fit in with the new environment they are going to be used in. There is no one approach that works for all cases as is. Due to this, there is a considerable amount of research that is done and is being done in further refining and discovering methods for automatic timetabling.

1.1 The University Course Timetabling Problem

The University Course Timetabling (UCT) Problem is a problem where l lectures are being placed into a t timeslots and r rooms in such a way that it will result in a feasible timetable [16]. It was proven by Cooper, T. and Kingston, J. that this problem is an NP-complete problem which means it is difficult to produce feasible timetables [10]. Identifying whether a timetable is feasible or not is determined by two types of constraints: (a) hard constraints, and (b) soft constraints

Hard constraints are constraints that are not supposed to be violated. Violation of these constraints immediately leads infeasibility of a potential. Some proposed approaches immediately discard timetables that are infeasible [16][17], while others would give high penalty costs for each hard constraint violated [13]. The set hard constraints may differ from one approach to another. However, a common hard constraint that can be observed across approaches is that a student, and, naturally, a teacher as well, cannot attend two or more classes in the same timeslot [16][4]. This is obvious as a person cannot be in two or more places at the same time. These hard constraints are what all constructed timetables must satisfy to be considered as a potential solution.

The counterpart to hard constraints are soft constraints. Unlike, hard constraints, soft constraints are *technically* optional for timetables to satisfy. However, finding the best timetables for a given problem require minimization of the number of soft constraints being violated [16]. The set soft constraints differ from one case to another. Another case might require soft constraints that are not present in another case. For example, one approach [17] has a soft constraint where students should not have only one class in a single day, while another [13] does not and instead does not allow

having more than three lectures of the same class schedule adjacently on the same day. Many approaches do not consider hard constraints in computing the fitness value of a timetable. Instead, many only use the soft constraints as part of their fitness function [18][17][4].

1.2 Algorithms Used In Approaches For Solving the University Course Timetabling Problem

Great Deluge

Many previous works dealing with course timetabling utilize an optimizing heuristic (or derivatives of it) called **Great Deluge**. Great Deluge was introduced by Gunter Dueck, and is a heuristic that is similar to Hill Climbing and Simulated Annealing. To understand how it works, imagine that you are in a point in some area with mountainous terrain. This area constitutes your solution space, with higher points in the area having higher values and, thus, having a better solution. The initial point you are located in in the area represents the initial solution that is generated. Imagine as well that it is raining endlessly and the water level W is continuously rising at a constant rate R_w , where $R_w > 0$. W can start at any value that is greater than 0. Assuming that we are attempting to maximize some function Q , which evaluates a solution based on some criteria, your goal is to locate the relatively highest point in the location. This point can be thought of as the local optimum. Locating the highest point involves walking around the area that is not below the current water level. This will force you to walk to a higher and higher point since W is constantly rising. Once you are no longer able to proceed to a higher level, it means that you are now in a local optimum. Going outside the analogy and back to a technical perspective, this

local optimum would now be the *relatively* best solution for your problem. Every "walk" or move to a higher point is nuanced relative to the analogy. A single walk means construction of a new solution S_{new} with basis on the current solution $S_{current}$. If $Q(S_{new}) \geq W$ [7], then we accept S_{new} as the new current solution and "walk" towards it, and we increase W by R_w . Otherwise, we simply generate a new S_{new} . These walks are performed until $Q(S_{new})$ is not greater than $Q(S_{curr})$ for a long time or we have reached the maximum number of moves/iterations [11]. Great Deluge can also be adapted to minimize Q . Instead of W increasing, it will be decreasing by the same rate. A solution S will now be accepted if $Q(S) \leq W$. Conversely, the algorithm will stop when $Q(S_{new})$ is not lesser than $Q(S_{curr})$. The second stopping condition for the algorithm still applies in this case [7][16][17]. This minimization case is the one adapted by Great Deluge-based works that focus on course timetabling.

Genetic Algorithms

Great Deluge is a metaheuristic [11] that has been used for solving course timetabling problems [7][16][17]. Aside from Great Deluge, another metaheuristic that has seen use in the course timetabling problem and its variations is the Genetic Algorithm. The genetic algorithm is an optimization algorithm whose behaviour is based on how nature works, particularly on how reproduction works at a genetic level. Many implementations of the genetic algorithm for course timetabling represents the problem by having each gene in the genetic representation of the timetabling be a two-element tuple which consists of the class and the agents that will partake in that said class [3][19][13]. Some use another representation. In the algorithm, an initial population is first generated. This population does not necessarily contain the locally optimal

solution but it is where the relatively best solution will be obtained from. This initial population is referred to as the first generation. From this population, a certain number of individuals will be randomly selected to be bred with one another and be the parents of the next generation. Selection of individuals is dependent on an individual's fitness, with the most fit individuals usually being selected. Calculating this fitness is dependent on the problem. In the context of course timetabling, fitness is based on the constraints that have been violated by the current solution/timetable generated [3][19][14][15][13][9][18]. The breeding process involves selection of parents and having genes from the parents crossover and/or mutate to produce new offspring. Crossover is when genes from both parents are combined to produce an offspring. On the other hand, mutation is done by changing a random gene from either parent to create an offspring. Determining which genes from either parent to apply onto the offspring and which to mutate is dependent on implementation. Once a population of new generations is established, the cycle repeats. This continuous reproduction of generations eventually produces solutions "moves" towards the optimal solution [2]. Despite being able to produce feasible solutions for university timetabling, it should be noted that using a genetic algorithm approach may require more time executing compared to other approaches due to its population-based property. When compared to simulated annealing, another approach for university timetabling, the genetic algorithm takes more time executing [14]. However, approaches utilizing the genetic algorithm can see an improvement in execution time when they utilize graphics processing units (GPUs). The work of Yousef, A., Salama, C., Jad, M., El-Gafy, T., Matar, M., and Habashi, S. showed that it is possible to speed up genetic algorithms using GPUs. In their work, they accelerated the computation of the fitness function.

Their experiments show that execution speed can be improved by up to 59 times in very large problem instances and by 280% overall when utilizing the GPU [21].

Chapter 2

Review of Related Literature

University course timetabling is a well-known problem. Many approaches has already been proposed before. These proposals usually use optimization-based algorithms. However, some approaches utilize some form of machine learning, specifically reinforcement learning, as part of their methodology. No single approach works best for all, and these approaches have their own tradeoffs.

2.1 Great Deluge-Based Works

One of the earliest works that uses Great Deluge was that of Burke, E., Bykov, Y., Newall, J., and Petrovic, S. [7]. Their work explores the use of Great Deluge in university course timetabling. In their work, they defined Q as the sum of the number of soft constraints violated. They also, however, slightly modified Great Deluge. We will be referring to the modified version as the Extended Great Deluge (EGD). Instead of simply obtaining a single new solution S in an iteration and comparing $Q(S)$ to W . Burke, E., et. al. extended Great Deluge to take multiple neighbouring solutions N on every iteration. A random solution S_N from N will be taken and it will be the one to be compared to W and the current solution. S_N will be accepted as the current

solution if $Q(S_N) \leq Q(S_{curr})$. S_N can still be accepted only under the condition that $Q(S_N) \leq W$. In addition to taking multiple solutions in an iteration, a few more extensions have been added to Great Deluge by Burke, et. al.. The initial value for W is set to $Q(S_i)$, where S_i is the initial solution. Computing for R_w is made easier by using the following equation:

$$R_w = \frac{W - Q(S')}{N_{moves}}$$

N_{moves} denotes the number of moves the algorithm will perform before terminating. The desired number of violations the desired solution S' should have is denoted by $Q(S')$. Through their work, Burke, E., et. al. observed that there is a tradeoff between the quality of produced timetables and the amount of searching time given to the algorithm. The authors noted that even though in some situations, a user would require obtaining the results quickly, it is naturally more preferable that one gives the algorithm more in order for it provide high quality timetables. As the authors have mentioned, course timetabling normally only occurs once or twice in a year. Thus, a long amount of time for searching "seems to be quite acceptable". EGD was compared to Simulated Annealing (SA), Threshold Acceptance (TA), Hill-Climbing (HC), and 21 algorithms submitted for the International Timetabling Competition (ITC) of 2002. The data set used for comparing the algorithms was the one provided by the aforementioned competition. The experiments of Burke, E., et. al. show that EGD produce less scattered results than SA, TA, and HC. This proves the effectiveness of EGD as a heuristic for course timetabling. Further proof of this effectiveness is shown by comparing the heuristic to the 21 other algorithms participating in the ITC. EGD generated the best results for 8 out of 23 data sets among the participating

algorithms [7].

Building upon the work of Burke, et. al. is that of Landa-Silva, D., and Obit, J.'s [16]. The primary contribution of Landa-Silva, D., and Obit, J. is the introduction of a modification of the Extended Great Deluge by Burke, E., et. al., the Non-Linear Great Deluge. Key to this heuristic is the use of an equation, instead of a constant rate, in determining the amount of reduction in the water level in an iteration. We refer to this amount as the decay rate. The next water level W is computed with the following equation:

$$W = W \times (\exp^{-\delta(\text{rand}[\text{min}, \text{max}])}) + P$$

P is "the minimum expected penalty corresponding to the best solution". Another key part of the equation is $\exp^{-\delta(\text{rand}[\text{min}, \text{max}])}$, which controls the speed the water level decreases. Besides from contributing an equation-based decay rate for the water level, another modification they added to the Extended Great Deluge is the conditional increase of the water level. If the current solution obtained S_{curr} is about to converge with the water level W , i.e. when $W - Q(S_{curr}) < 1$, then the Non-Linear Great Deluge algorithm allows the water level to rise for a certain amount. This amount is a random value from the interval $[W_{min}, W_{max}]$. The intended increase is to allow the algorithm to "accept slightly worse solutions to explore different areas of the search space in the hopes of finding better solutions". The work of Landa-Silva, D. and Obit, J. uses three moves in generating solutions:

- **M1:** Selects one random class and gives it a random but feasible timeslot-room pair.

- **M2**: Selects two random classes and swaps their timeslot-room pairs while maintaining feasibility.
- **M3**: Looks for a class which violates soft constraints based on their current timeslot-room pair. It then moves this class to a random timeslot-room pair but still maintaining feasibility.

It should be noted that these three neighbourhood moves always maintain compliance with the hard constraints. Another key modification the authors introduced is the three-step heuristic for generating an initial solution. This heuristic is based on the work by Chiarandini, M., Birattari, M., Socha, K., and Rossi-Doria, O. [8]. The three steps are as follows:

1. All classes will first be assigned a random timeslot. However, unassigned classes with the highest number of conflicts will be assigned first. A class has a conflict with another class if it has students also taking the other class. The maximum bipartite matching algorithm [1] is then used to assign each event to a room. This first step of creating an initial timetable S does not guarantee feasibility. But the solution will be further improved by the later steps.
2. Moves M1 and M2 are then used to improve the current S . At this step, feasibility and satisfaction of the hard constraints are sought. As such, "a move is only accepted if it improves the satisfaction of the hard constraints". This step performed continuously until there are no more improvements to S after 10 iterations.
3. Tabu search [6] is then used to further refine S . In this step, classes that were assigned t_{iter} iterations ago will be stored in the tabu list. t_{iter} is computed as

$t_{iter} = ran(10) + \delta \times N_v$, where $ran(10)$ is a random number from the interval $(0, 10)$, $\delta = 0.6$, and N_v is the number of classes that partook in violating the hard constraints. The termination condition for this step is when after 500 iterations, no solution has been produced that is better than the current best.

Only step 1 is run once. The other steps, steps 2 and 3, are performed repeatedly in order until a feasible solution/timetable is obtained. In all 11 experiment instances obtained from the work of Socha, K., Knowles, J., and Sampels, M. [20], which has 5 small instances, 5 medium instances, and 1 large instance, Non-Linear Great Deluge produced results better than Great Deluge. Interestingly enough, in 4 of these experiment cases, NLGD performed the best compared to the best known literature at the time NLGD was introduced [16].

2.2 Machine Learning-Based Works

All related works encountered that use machine learning that were proposed do not use a purely machine-based approach. Rather, they combine heuristics with machine learning.

An example of such works, is that of Obit, J., Landa-Silva, D., Sevaux, M., and Ouelhadj, D. [17]. Their work is built upon their previous work on Non-Linear Great Deluge for course timetabling [16]. An addition to their work compared to previous studies is the use of reinforcement learning as part of their solution generation process. Previous works, relative to their's, randomly select moves. But, Obit, J., et. al. uses reinforcement learning to determine which moves is best to further refine the current solution. Two types of reinforcement learning are employed and investigated in their work: **(a)** reinforcement learning with static memory length, and **(b)** reinforcement

learning with dynamic memory length. Initially, all moves are given equal weights of 0.01, with the weight for each move i denoted by w_i and thus, equal probabilities of being chosen for an iteration. The probability for each move is denoted by p_i is computed by

$$p_i = \frac{w_i}{\sum_{i=1}^n w_i}$$

n is the number of moves there are. In the type with static memory length, on every iteration, the moves are punished or rewarded depending on their performance. A move is rewarded by giving its weight 1 point. It is punished by giving it no points at all. The weights are updated every learning period lp computed using $lp = \max(N_{moves}/500, n)$, where the total number of feasible moves performed is denoted by N_{moves} . During this learning period, the water level is also increased by a certain amount. Performance of each move is computed based on the number of it was called, number of times it generates solutions that has different fitness values, and the number of times it produces solutions that have been accepted. On the other hand, reinforcement learning (RL) with dynamic memory length, the probability p_i of a move being chosen is computed through

$$p_i = \frac{w_i + w_{min}}{\sum_{i=1}^n w_i + w_{min}}$$

In this equation, $w_{min} = \min(0, w_i)$. Unlike its counterpart, this type of RL updates the weights of the moves every time a feasible move is performed. The performed move is rewarded when it produces an improve solution, and punished otherwise. Differing from its counterpart once again, this type of RL uses a piecewise function R to compute the reward/punishment for the currently selected move i at

the current iteration j . In the function, Δ is the difference between the best solution so far and the current generated solution.

$$R = \begin{cases} 1 & \text{if } \Delta < 0 \\ -1 & \text{if } \Delta > 0 \\ 0.1 & \text{if } \Delta = 0 \text{ and new solution} \\ -0.1 & \text{if } \Delta = 0 \text{ and no new solution} \end{cases}$$

Each weight w_i is then computed at each iteration h with the formula below

$$w_i = \sum_{j=n_{timeslots}}^h \sigma^j R$$

$n_{timeslots}$ is the number of timeslots there are. The parameter σ is a random value between $(0.5, 1.0]$. The parameter's value is set every learning period lp . Similar to the previous type of RL, the lp is still determined using the same formula, and the water level increases on every learning period. The experiments of this work included an experiment comparing the two types of reinforcement learning employed. Both types of RL produced optimal results in the small problem instances and good results in the medium instances. In the large problem instance, the static memory RL produced a better result than its counterpart. When comparing to other Great Deluge-based works, namely EGD, NLGD, ENLGD, and GD, the static memory RL produced the best solution in all problem instances except the large instance, in which its solution's score has a difference of only 1 from the best solution, which was generated by the Extended Great Deluge. The dynamic memory RL produced the worst solution among the algorithms compared for the large problem, but came in second overall in all of the medium problem instances, but one in which it came in third. When comparing the algorithm to hyper-heuristic-based algorithms with the same problem

instances, the algorithm, specifically the static memory RL-based algorithm, produced the best results, except for the large instance. The authors also noted that the value of lp affects the quality of the best solution the algorithm produces, with the sweet spot being either $lp = 2500$ or $lp = 5000$. When the algorithm is finally compared to course timetabling algorithms that are known to produce the best results for a specific problem instance, the algorithm, specifically the static memory RL one, produced new best solutions in all medium problem instances. This proves that the algorithm is a viable solution for course timetabling.

Aside from reinforcement learning, neural networks, specifically Hopfield neural networks, have also been used in timetabling. Lindblad, M. compared a timetabling algorithm that uses Hopfield neural networks with that of heuristics based approaches, namely simulated annealing

2.3 Genetic Algorithm-Based Works

A work that utilizes genetic algorithms is the work of Alves, S., Oliveira, S., and Rocha Neto, A. [3]. Unlike the previous works which tackles the problems as organizing a single set of lectures onto a table of time slots (which can be thought of timetabling for a semester), Alves. S., et. al.'s work organizes a timetable for entire courses throughout multiple semesters. We refer to a course here as a degree a student takes up in college. It is important to note that the perspective in which they tackle the problem is that of the Brazilian university academia's. Unique to their approach is that they only consider a single course in timetabling at a time. This means that the classes of other courses are not **directly** taken into account. Instead, what is considered is the unavailability of agents, which are students and professors. This

unavailability determined from the assignments of the agents from the timetabling of the previous courses. Each course will contain timetables for each semester. In constructing the timetable for a course, a genetic algorithm is applied on each course, with each individual being a timetable for the course. Specific rules for selection of parents have not been specified in the paper but it can be presumed that the fittest individuals are chosen for breeding. Fitness F is computed using the following function:

$$F(C) = 1 - \frac{AM_C + AU_C + AL_C}{AM_{wc} + AU_{wc} + AL_{wc}}$$

C is a course timetable. AM represents the number of times an agent has been assigned to concurrent events. AU stands for the number of times agents have assigned to classes whose timeslots they are not available in. Lastly, AL represents the number of times a class has been given more than three consecutive timeslots. The numerator represents the values for C , while the denominator denotes the worst case values of those variables. Crossover is performed using the OX operator from the work of Chinnasri, W., Krootjohn, S., and Sureerattanan, N. [9]. Mutation is performed in a course by selecting a random semester timetable and swapping two random timeslots. The timetable of a course is then stored in a global timetable. This timetabling is performed until all courses have been timetabled. The final solution is the global timetable. According to the experimental results, the approach of Alves, S., et. al. produces timetables that have a fitness close to 1. However, producing fit timetables required performing the approach multiple times with different values. In their work, the authors obtained the the parameter values by performing 15 tests with each test having different values for the parameters. It should be highlighted

that the stop criteria has been set to when a fitness of 1 has been achieved or the execution time has reached 10 minutes. Alves, S., et. al. found that, in their problem context, the parameters that produced the best results are: (a) 25% mutation rate, (b) 50% OX crossover rate, and (c) a population size of 75. Using these parameters, the authors managed to have their approach produce 493 generations and having a run time of just 3.7 minutes, with all timetables for each course having a fitness of approximately 1.

Innet, S. has also utilized a genetic algorithm for timetabling in Thai universities, at least for the University of the Thai Chamber of Commerce [13]. Unlike the other aforementioned works, the work of Innet, S. focuses on examination timetabling, instead of course timetabling. However, his approach may still be utilized for course timetabling as the former is similar with the latter. The crossover operation that was chosen for Innet, S.'s work is a simple operation of selecting random genes from one parent and placing them onto a child chromosome in the same position as they were in the parent. This would entail that unfilled slots will be present. These slots will be filled by all genes, except those that have already been selected from the first parent, from the other parent with the first gene filling the first unfilled slot, the second with the second slot, and so on. The mutation operation is simpler. A parent will be selected and its genes will be copied onto a child. Two random genes will be selected and swapped. Crossover and mutation are not performed consecutively. Rather, one of them will be performed depending on the crossover probability parameter. During breeding for the next generation, only two individuals will be selected as the parent with the condition that these parents have the best fitness value. Based on the experimental, Innet, S.'s approach manages to produce fit timetables (only having a penalty cost of

32) but only when the crossover probability is set to 75%. Unfortunately, his approach was not compared to other approaches from known literature. As such, it cannot be determined how well his approach performs compared to other approaches such as Great Deluge in examination timetabling, especially in the case he is applying his proposed approach [13].

Another genetic algorithm-based approach is that of Bedoya, C. F., and Santos, M. [4]. A key point to take note with the approach of the authors is that, unlike previous works (especially the aforementioned ones), the crossover operation standard in genetic algorithms is not used. The reasoning behind this is that the operation can introduce violations to the hard constraints. This will necessitate the repairing any resulting timetables to satisfy the violated constraints. As such, only the mutation operation has been used. It was noted by the authors that a timetable already intrinsically has all the necessary information to produce a feasible timetable. Properly sorting this timetable is all that is required to produce a good timetable. Mixing the genes of two would-be parents will not improve the solution. The mutation operation proposed by Bedoya, C. F., and Santos, M. is not completely based upon randomness. Instead, a set of rows will be swapped with another set of rows. These rows are determined by two swapping points, where the slicing of the chromosome for swapping will start, and the depth parameter which determines how many rows will be swapping. Two random genes from the chromosome are selected as the swapping points for the mutation operation. The probability in which a co-beliigent gene will be chosen as a swapping point is determined by a parameter. The fitness function F in the work of Bedoya, C. F., and Santos, M. is defined as:

$$F_{obj} = \sum w_i R_i$$

where R_i represents the number of times a constraint was violated and is weighted by a value w_i , which represents the priority of the constraint. It must be noted that the authors' work seeks to maximize this fitness function. Each of the unsatisfied constraints gives a negative value, making the fitness for the optimal solution to be 0. The experiments of Bedoya, C. F., and Santos, M. show that their non-standard genetic algorithm is capable of producing optimal timetables. It was also observed that increasing the population size reduces the number of iterations needed to produce the optimal solution. However, further increasing the population size beyond 20 gives back diminishing returns especially on system resources expended. The number of iterations needed to obtain the optimal solutions stagnates and remains the same at around 250 when going beyond that threshold point of 20. However, in the same vein as the work of Innet, S. [13], the performance of the non-standard genetic algorithm cannot be determined since the authors did not conduct experiments to compare their approach against previous approaches for course timetabling [4].

Course timetabling using a genetic algorithm procedure has also been performed in a non-university setting. Raghavjee, R. and Pillay, N. applied the genetic algorithm for a primary and high school in South Africa [19]. Similar to the work of Bedoya, C. F., and Santos, M., the authors discarded the crossover operation and only used the mutation operation during the breeding process. Their proposed approach has two phases: The first phase focuses on produces feasible timetables, and the second phase will focus on improving the produced timetables. The two phases all utilize genetic algorithms to accomplish their tasks. The timetable representation used in

this work is a 2D matrix where the rows are the time slots and the columns being the classes. The intersections of the rows and classes are where the teachers are placed in. Generation of the initial population involves creating a m -size population. The fittest individuals from that population are gathered to create the initial population. The fitness function used is dependent on which phase the fitness is being computed. In the first phase, the hard constraint cost of the timetable is used. On the other hand, the soft constraint cost is used during the second phase. During timetabling, the tuple of class and teacher are sorted based on the difficulty of scheduling the said tuple. This difficulty is determined by a set of low-level heuristics. Each tuple will be given a time slot that gives the tuple the least penalty as possible. The first phase will have this initial population of timetables be refined for a number of generations. During selection of parents for the next generation, a variation of the tournament selection. At the end of this phase, feasible timetables have already been generated. The last generation will now then be used by the second phase. Two mutations operations were used: (a) simple mutation, and (b) a hill-climber operator. The simple mutation simply swaps two random tuples in the same class. The second operator simply performs swaps two teachers and checks if the swapping improves the fitness of the timetable. If it does, then the swap is accepted. Otherwise, more swapping is performed until an improvement occurs. However, there is a limit to the number of times attempts at swapping is performed. This limit is set by a parameter. The experiments of the authors show that a feasible timetable is possible with the proposed method. The minimum hard constraint cost mustered by all the generated timetables is 0 for both primary and high school, with the average hard constraint cost at 1.1 and 2, respectively. Only the high school timetable has soft

constraint violations with the minimum cost at 2 and the average cost at 2.67. Unlike the previous two discussed works, the authors compared their approach to another previous work, specifically the approach of Beligiannis, G. N., Moschopoulos, G. P., Kaperonis, P., and Likothanassis, S. D. [5]. The minimum hard constraint cost of each approach is 0. However, the minimum soft constraint cost of Raghavjee, R. and Pillay, N. is just 2, compared to the cost of 6 of the approach they are comparing against. Despite having better results than a previous work, it would have been more beneficial to compare their approach to other different approaches from previous works for us have better insights into the proposed approach [19].

Genetic algorithms may also be combined with another approach in producing timetables, just like how reinforcement learning was integrated into great deluge in a previously discussed work. Feng, X., Lee, Y., and Moon, I. formulated a new approach combines integer programming and genetic algorithms together to produce feasible university timetables [12]. Feng, X., et. al.

Chapter 3

Statement of the Problem

Raster images are composed of a pixel matrix. This makes their data representation simple. However, this reduces the amount of detail and quality they have. This limitation is clearer when scaling raster images. This motivates the use of an alternative form of representing images. One form is vector images, which uses mathematical equations to represent an image. The process of converting a raster image to a vector image is called *vectorization*. Semi-structured imagery, such as those used in graphic designs, is one of the classes of images that would benefit from vectorization. This process will allow such images to be easily scaled without sacrificing quality nor detail.

There have been numerous works that tackle image vectorization for semi-structured images. Many of these works primarily use a curve optimization algorithm such as variants of NEWUOA and conjugate gradient. A machine learning approach has been used in one of the works, but as a preprocessing step only [?]. Deep learning have been used to solve problems in multiple domains such as natural language processing (NLP), object detection, and playing board games. No other known work has applied deep learning as a core step in image vectorization. This study will deal with such application.

Chapter 4

Objectives

This study is primarily aims to apply deep learning via artificial neural networks to the problem of vectorizing semi-structured imagery. However, there are still some key objectives that this study seeks to accomplish:

1. To develop an approach that considers Gestalt psychology.
2. To evaluate the effectiveness of using deep learning for image vectorization.
3. To evaluate the accuracy of results obtained from the proposed approach to the target raster inputs.
4. To evaluate and compare the results of the proposed approach to that of previous semi-structured image vectorization methods.
5. To evaluate and compare the speed of the proposed approach compared to previous semi-structured image vectorization methods.

Chapter 5

Proposed Methodology

The proposed methodology will be based on the frameworks proposed by Hoshyari, S., et. al. [?], Yang, M., et. al. [?], Xiong, X., et. al. [?], and Laube, P., et. al. [?]. Additionally, human perception will also be taken account. Thus, the Gestalt psychology principles of accuracy, simplicity, continuity, and closure, as taken from Hoshyari, S., et. al., will be taken into account as well.

5.1 Image Preprocessing

Before a raster input image can be fitted with curves, it must preprocessed to simplify the vectorization procedure and align it with .

Corner Detection

The first step in the approach is detecting the corners in the input image. This is an important step as it will allow us to enforce the simplicity principle in Gestalt psychology and make sure the resulting vectorization be C^0 continuous should the raster input be as such.

This stage will be based from the corner detection classifier of the work by Hoshyari, S., et. al. [?]. A random forest classifier is used in the said work. Supervised learning is used as corners are manually annotated. Annotated corners will not be specific pixels. Rather, corners will be between at least two pixels. Training data is available publicly provided by the researchers. However, such data is limited only to quantized data (i.e. aliased data). The target raster input is expected to be anti-aliased data. As such, the training data will have to be built from scratch to support anti-aliased data.

Region Segmentation

In line again with the simplicity principle, the input image must be divided into regions. This will result in simpler curves being used in the final vectorization output. The additional benefit of segmenting the input into multiple distinct regions is the possibility for parallelism to be used in the vectorization approach. Since each region is distinct and independent from one another, multiple regions can be vectorized at the same time. Thus, speeding up the vectorization process.

Due to the nature of semi-structured imagery where each region will only contain a single colour, a scanline-based approach can be used in this stage. The scanline algorithm will be based off of the scanline algorithm used for boundary pixel detection in the work of Xiong, X., et. al..

For each line l_i , where $0 \leq i < h$ | h is the height of input image, in the raster input, each pixel p_i will be assigned to a pixel set r_{ij} , where j is the index to a set in l_i . Each pixel set will contain horizontally adjacent pixels that have the same or near-same colours. We include pixels whose colours are within a certain threshold k from the colour of the pixels that is most prominent in the set. This threshold is

necessary due to the fact that certain parts of a regions may contain an anti-aliased pixel. Each l_i will have a pixel set vector r_i containing all pixel sets of l_i :

$$r_i = (r_{i0}, r_{i1}, \dots, r_{i(j-1)}, r_{ij})$$

Consequently, a region vector r will contain all pixel set vectors.

$$r = (r_0, \dots, r_i)$$

Note that there is an opportunity to utilize parallelism, as shown in the work of Xiong, X., et. al. [?], during this step due to the independent nature of every line in the raster input.

Once we obtain all the pixel sets for each line, we iterate through r , $(h - 2)$ times. For each iteration, we process r_i and r_{i+1} . If there are any r_{ij_α} whose pixels are vertically adjacent and have the same colour (or within k) to another pixel set $r_{(i+1)j_\beta}$, then those two pixel sets are merged into one. By the end of the iterations, we have obtained a set of regions that we can individually vectorize.

5.2 Curve Approximation

The core step of the proposed approach is curve approximation. This stage fits curves to the region boundaries of the raster input. This stage is based on the work by Laube, P., et. al. on curve approximation on point sequences using deep learning [?].

Point Sequence Generation

The work of Laube, P., et. al. takes a point sequence as input. As such, for this proposed approach, we must generate point sequences from the region we will be

vectorizing. For every region we ought to vectorize, we treat the center of boundary pixels and the detected corners of each region as a point sequence. However, we must also take into account the fact that certain portions of a region may be a corner where the curves in such segment would have C_0 continuity. As such, the point sequence we generate must take into account corners. This would implore us to take note of the following cases during point sequence generation:

1. For regions with no corners, a random pixel will be selected as both start and end point of the point sequence. The expectation is that there will be no difference in the resulting curve from choosing a different start and end point.
2. For regions with a single corner, the corner point will be selected as the start and end point of the point sequence. This is to ensure that the resulting vector output will have a corner at that point.
3. For regions with two or more corners and assuming n is the number of corners, the point sequence will be divided at those corners into separate point sequences. This will result in $(n + 1)$ new point sequences. Each new point sequence will have their start and end points be the corner points they are adjacent to.

Each point sequence will then be passed to the next stage to be parametrized and have a curve approximated for.

Point Sequence Curve Approximation

The point sequences obtained from the previous step are now to be fitted with curves, specifically B-splines. As provided by the framework by Laube, P., et. al., two neural networks will be used in this stage and some preprocessing will be performed on

the point sequences. The two neural networks are a point parametrization network (PPN), which approximates parametric values to point sequences, and a knot selection network (KSN), which predicts new knot values for knot vector refinement.

Sequence Segmentation

The input point sequence must first be split. This is to ensure that real data and training data match in terms of complexity. Let us define a function $\hat{k}(p)$ that measures the complexity of a point sequence p , where k_i is the curvature at point p_i , given its total curvature:

$$\hat{k}(p) = \sum_{i=0}^{m-1} \frac{(|k_i| + |k_{i+1}|) \|p_{i+1} - p_i\|_2}{2}$$

A point sequence p is split into point sequence segments $p^s, s = 1, \dots, r$ at the median, if $k(\hat{p}) > \hat{k}_t$ for a threshold \hat{k}_t . This \hat{k}_t will be set, as per the original authors have done, to the 98th percentile of $\hat{k}(\cdot)$ of the training set. This process is performed $r - 1$ times, until each p^s satisfies $k(\hat{p}) > \hat{k}_t$.

Sub/Supersampling and Normalization

To be able to approximate parametric and knot values using the PPN and KSN, the number of points per segment p^s must equal the input size l of the aforementioned networks. As such, all segments p^s are either subsampled or supersampled.

If a segment p^s has a number of points greater than l , then p^s is subsampled. This process involves drawing points in p^s such that the drawn indices i are equally distributed and include the first and last point. If, on the other hand, the number of points in p^s is less than l , then *temporary* points are linearly interpolated between consecutive points p_i^s and p_{i+1}^s . This interpolation is performed until the number of

points equal l .

The sampled segments are then normalized to \bar{p}^s , which consists of the points

$$\bar{p}_i^s = \frac{p_i^s - \min(p^s)}{\max(p^s) - \min(p^s)}$$

where $\min(p^s)$ and $\max(p^s)$ are the minimum and maximum coordinates of p^s respectively.

Parametrization of Point Segments

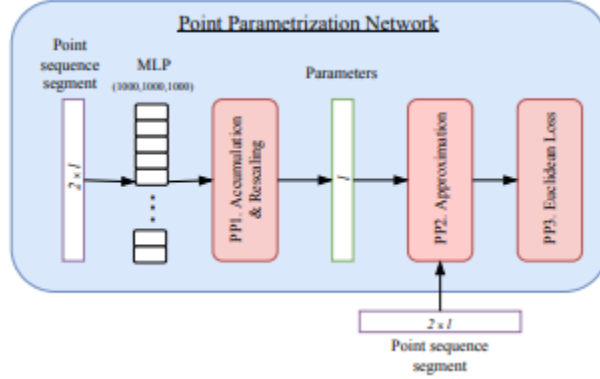


Figure 5.1: The network architecture of the Point Parametrization Network (PPN)

The PPN will be responsible for parametrization of point segments. For every \bar{p}^s , the PPN generates a parametrization $\bar{t}^s \subset [0, 1]$, which will be rescaled to $[u_{s-1}, u_s]$ and adapted to sampling of \bar{p}^s .

In supersampled segments, the parameters t_i^s of temporary points are simply removed from \bar{t}^s . In a subsampled p^s , for every point p_i that was removed from the segment, a parameter \bar{t}_i is inserted to \bar{t}^s .

$$t_i = t_\alpha^s + (t_\beta^s) \frac{\text{chordlen}(p_\alpha^s, p_i)}{\text{chordlen}(p_\alpha^s, p_\beta^s)}$$

$chordlen$ is the length of the polygon defined by a point sequence. In the subsampled segment, with parameters t_α^s and t_β^s , p_α^s and p_β^s are the closest neighbours of p_i .

The initialization of the parametric step requires an initial knot vector. We first define $u_0 = 0$ and $u_n = 1$. For each segment (except the last one), one knot u_i is added.

$$u_i = u_{i-1} + \frac{chordlen(p^s)}{chordlen(p)}, i = 1, \dots, r - 1$$

This yields a start and end knot for every point sequence segment.

The PPN Architecture The PPN, as stated earlier, takes in an input of segments p , which can be written as $p = (x_0, \dots, x_{l-1}, y_0, \dots, y_{l-1})$. The parameter domain is defined as $u_0 = t_0 = 0$ and $u_n = t_{l-1} = 1$. For a sequence of points p , a parameter vector $t = (t_i)_i$, is defined as $t_i = t_{i-1} + \Delta_{i-1}$. The task of the PPN is to predict missing values $\Delta = (\Delta_0, \dots, \Delta_{l-2})$ with

$$\Delta_{subi} > 0, i = 0, \dots, l - 2$$

such that $t_0 < t_1$ and $t_{l-2} < t_{l-1}$. We apply a multilayer perceptron (MLP) to the input data p , yielding as output a distribution for parametrization $\Delta^{mlp} = (\Delta_0^{mlp}, \dots, \Delta_{l-2}^{mlp})$ of size $l - 1$.

The PPN further contains additional layers PP1, PP2, and PP3, which will be discussed next.

PP1. Accumulation and Rescaling The output Δ^{mlp} is used to compute a parameter vector t^{mlp} with $t_0^{mlp} = 0$ and

$$t_i^{mlp} = \sum_{j=0}^{i-1} \Delta_j^{mlp}, i = 1, \dots, l-1$$

Since t_{l-1}^{mlp} is usually not 1, rescaling t^{mlp} yields the final parameter vector t with

$$t_i = \frac{t_i^{mlp}}{\max(t^{mlp})}$$

The MLP layer in the PPN uses a softplus activation function defined to be:

$$f(x) = \ln(1 + e^x)$$

PP2. Approximation B-spline curve approximation is included directly into the PPN as a network layer. The input points p and their parameters t are used for an approximation with knot vector $u = (0, 0, 0, 0, 1, 1, 1, 1)$ for $k = 3$. The approximation layer's output $p^{app} = (p_0^{app}, \dots, p_{l-1}^{app})$ is the approximating B-spline curve evaluated at t .

PP3. Euclidean Loss A loss function, which is a Euclidean loss function, is to be used in the PPN. The Euclidean Loss function is defined as

$$\frac{1}{l} \sum_{i=0}^{l-1} \|p_i - p_i^{app}\|_2 \quad (5.2.1)$$

Parametrization Refinement

In some cases, the approximated parametrization of a p^s may have errors. The approximation error of a p^s is computed by using the Hausdorff distance to the input data p .




Figure 5.2: The network architecture of the Knot Selection Network (PPN)

The segments p^s that have a large approximation error are computed a new knot using the KSN. For \bar{p}^s and \bar{t}^s , the KSN generates a new estimated knot $\bar{u}_s \in [0, 1]$. The new knot \bar{u}_s is mapped to the actual knot value range $[u_{s-1}, u_s]$ by

$$\tilde{u}_s = u_{s-1} + \bar{u}_s(u_s - u_{s-1})$$

Instead of \tilde{u}_s , the parameter value t_i closest to \tilde{u}_s is inserted into u . u can be further refined until the desired curve approximation error threshold is satisfied.

The KSN Architecture This stage utilizes a KSN, as mentioned earlier. The KSN predicts a new knot u to the interval $(0, 1)$ for a given segment p and parameters t , of which were previously approximated by the PPN. This network uses an MLP, which transforms the input to a single output value u^{mlp} . The RELU function is used as the activation function for the MLP, except for the output layer where the Sigmoid function is used instead.

Similar to that of the PPN, the KSN has three additional layers: KS1, KS2, and KS3.

KS1. Threshold Layer The new knot u computed has to satisfy the following conditions: $u \in (0, 1)$, $t \cap [0, u] \neq \emptyset$, and $t \cap [u, 1] \neq \emptyset$. Satisfying these conditions will require us to use a threshold layer which maps u^{mlp} to

$$u = \begin{cases} \epsilon & , \text{ if } u^{mlp} \leq 0 \\ 1 - \epsilon & , \text{ if } u^{mlp} \geq 1 \\ u^{mlp} & , \text{ otherwise} \end{cases}$$

Introducing a small $\epsilon = 1e - 5$ makes sure that the knot multiplicity at the end knots stays equal to k .

KS2. Approximation Approximation in the KSN is generally similar to that of PPN. The only different is that of the knot vector. The knot vector in the KSN is defined to be $u = (0, 0, 0, 0, u, 1, 1, 1, 1)$. For backpropagation, the derivative of the B-Spline basis functions with respect to u is required.

KS3. Euclidean Loss The loss function for the KSN is the same as that of the PPN. See 5.2.1.

Network Training

The training of the PPN and KSN will be based from the work of Laube, P., et. al. [?]. The input size of the network will be $l = 100$.

The data set generated by the original authors consists of 150,000 curves. This data was synthesized from B-spline curves. Random control points c_i were generated using a normal distribution μ and variance δ to define cubic ($k = 3$) B-spline curves with $(k + 1)$ -fold end knots and no interior knots. The y -coordinates are given the configuration: $\delta = 2$ and $\mu = 10$. For the x -coordinates, $\delta = 1$ and $\mu = 10$ are used for

the first control point. All consecutive points have μ increased by $\Delta\mu = 1$. Curves with self-intersections are discarded, because the sequential order of their sampled points is not unique, and point sequences are usually split into subsets at the self-intersections. Smaller δ for the x-coordinates of control points reduces the number of curves with self-intersections. To closely match the target input as much as possible, we also include curves that have been manually fitted to raster images. These curves can be obtained from vector images available online.

For each curve, l points $p = (p_0, \dots, p_{l-1})$ are sampled. These curves then to have increasing x-coordinates from left to right. As such, index-flipped versions of the point sequences of the dataset are added, resulting in 300,000 point sequences. 20% of the sequences are used as test data in the training process.

The PPN is trained first since the KSN requires point parametrizations t , which is obtained from the PPN. After training, the PP2 and PP3 layers are discarded and PP1 becomes the output layer of the PPN. The parametric values t are computed for the training dataset by applying the PPN and train the KSN on the combined input. After training, KS2 and KS3 are discarded, with KS1 becoming the network output layer. The MLPs of the PPN and KSN will consist of three hidden layers with sizes (1000, 1000, 1000) and (500, 500, 500) respectively. Dropout is applied to the MLP layers. The network is trained using the Adam optimizer.

5.3 Region Colouring and Merging

Once the curves have been approximated, the vectorization of the region will be filled with the colour prominent in the raster version of the region. The regions will be plotted unto their locations in the original raster input. Once all the regions have

been plotted, they will be grouped into a single vectorization. This will now be the vectorization output of the raster image.

Chapter 6

Describing How You Validated Your Approach.

Chapter 7

Stating Your Results and Drawing Insights From Them.

Chapter 8

Summarizing Your Thesis and Drawing Your Conclusions.

Appendix A

What should be in the Appendix

What goes in the appendices? Any material which impedes the smooth development of your presentation, but which is important to justify the results of a thesis. Generally it is material that is of too nitty-gritty a level of detail for inclusion in the main body of the thesis, but which should be available for perusal by the examiners to convince them sufficiently. Examples include program listings, immense tables of data, lengthy mathematical proofs or derivations, etc.

Bibliography

- [1] Maximum bipartite matching - geeksforgeeks. *GeeksforGeeks*.
- [2] What is the genetic algorithm?- matlab & simulink.
- [3] Shara S. A. Alves, Saulo A. F. Oliveira, and Ajalmar R. Rocha Neto. A novel educational timetabling solution through recursive genetic algorithms. 01 2015.
- [4] Cristina Fernandez Bedoya and Matilde Santos. A non-standard genetic algorithm approach to solve constrained school timetabling problems. pages 26–37, 01 2003.
- [5] Grigorios N. Beligiannis, Charalampos N. Moschopoulos, Georgios P. Kaperonis, and Spiridon D. Likothanassis. Applying evolutionary computation to the school timetabling problem: The greek case. *Computers & Operations Research*, 35:1265–1280, 01 2008.
- [6] Jason Brownlee. Tabu search - clever algorithms: Nature-inspired programming recipes. *cleveralgorithms.com*.
- [7] Edmund Burke, Yuri Bykov, James Newall, and Sanja Petrovic. A time-predefined approach to course timetabling. *Yugosl. j. oper. res.*, 13:139–151, 01 2003.
- [8] Marco Chiarandini, Mauro Birattari, Krzysztof Socha, and Olivia Rossi-Doria. An effective hybrid algorithm for university course timetabling. *J Sched*, 9:403–432, 01 2006.

- [9] Wutthipong Chinnasri, Soradech Krootjohn, and Nidapan Sureerattanan. Performance study of genetic operators on university course timetabling problem. *IJACT*, 4:61–71, 01 2012.
- [10] Tim B. Cooper and Jeffrey H. Kingston. The complexity of timetable construction problems. pages 281–295, 01 1996.
- [11] Gunter Dueck. New optimization heuristics: The great deluge algorithm and the record-to-record travel. 104:86–92, 01 1993.
- [12] Xuehao Feng, Yuna Lee, and Ilkyeong Moon. An integer program and a hybrid genetic algorithm for the university timetabling problem. *Optimization Methods and Software*, 32:625–649, 01 2017.
- [13] Supachate Innet. A novel approach of genetic algorithm for solving examination timetabling problems: A case study of thai universities. 01 2013.
- [14] Johan Jonasson and Eric Norgren. Investigating a genetic algorithm- simulated annealing hybrid applied to university course timetabling problem: A comparative study between simulated annealing initialized with genetic algorithm, genetic algorithm and simulated annealing. 01 2016.
- [15] Kuan Yik Junn, Joe Henry Obit, and Rayner Alfred. The study of genetic algorithm approach for educational timetabling problems. 02 2018.
- [16] Dario Landa-Silva and Joe Henry Obit. Great deluge with non-linear decay rate for solving course timetabling problems. 01 2008.
- [17] Joe Henry Obit, Dario Landa-Silva, Marc Sevaux, and Djamila Ouelhadj. Non-linear great deluge with reinforcement learning for university course timetabling. 01 2011.

- [18] Sanjay R. and Rajan S. An application of genetic algorithm for university course timetabling problem. *IJAIS*, 11:26–30.
- [19] R Raghavjee and N Pillay. Using genetic algorithms to solve the south african school timetabling problem. 01 2010.
- [20] Krzysztof Socha, Joshua Knowles, and Michael Sampels. A max-min ant system for the university course timetabling problem. pages 1–13, 01 2002.
- [21] Ahmed H. Yousef, Cherif Salama, Mohammad Y. Jad, Tarek El-Gafy, Mona Matar, and Suzanne S. Habashi. A gpu based genetic algorithm solution for the timetabling problem. 01 2016.